

Advanced Network Security and Architectures

Project Report:

Kubernetes & Docker Swarm



Félix Saraiva – 98752

Table of Contents:

Table of Figures:.....	2
1. Kubernetes networking	4
Network Topology:	4
1.1. Kubernetes control plane.....	5
The Control Plane:.....	5
1.2. Deploying a single pod	7
Proof of Concept:	7
1.3. Deploying a single pod through a YAML manifest	9
Proof of Concept:	9
1.4. Deploying a ReplicaSet	10
Proof of Concept:	10
1.5. Deploying a ClusterIP service	13
Proof of Concept:	14
1.6. Kubernetes DNS	16
Proof of Concept:	16
1.7. Deploying a NodePort service	17
Proof of Concept:	17
2. Docker Swarm	20
2.1. Deploying a Service in Docker Swarm.....	20
Swarm mode Routing Mesh.....	20
Network Architecture	21
Proof of Concept:	21
References:	24
Annex A: Configurations	25
Docker and Kubernetes installation in GNS3:.....	25
Deploying a single pod	27
Deploying a ReplicaSet.....	28
Deploying a ClusterIP service.....	29
Kubernetes DNS	30
Deploying a Service in Docker Swarm	31

Table of Figures:

Figure 1 - Kubernetes networking network topology.....	4
Figure 2 - Kubernetes cluster components (Kubernetes, 2022).....	5
Figure 3 - Control plane packet from worker (UB1) to the master (UB2)	6
Figure 4 - Control plane packet from master (UB1) to the worker (UB2)	7
Figure 5 - mypod IP address and node	7
Figure 6 - Inside mypod shell	8
Figure 7 - GET request to the nginx server from mypod	8
Figure 8 - mypod YAML Manifest	9
Figure 9 - Pod creation and Information	9
Figure 10 - Repeating the creation command	10
Figure 11 - Modified YAML Manifest with manual node choice	10
Figure 12 - mypod created in node UB2	10
Figure 13 - ReplicaSet YAML Manifest.....	11
Figure 14 - ReplicaSet creation and Pods characteristics	11
Figure 15 - Pod's IP addresses and assigned nodes.....	11
Figure 16 - Pod deletion and replacement	12
Figure 17 - Scaling the ReplicaSet to seven Pods.....	12
Figure 18 - Deleting the entire ReplicaSet	12
Figure 19 - ClusterIP service diagram	13
Figure 20 - Replica and Curl Pods	14
Figure 21 - ClusterIP service description.....	14
Figure 22 - Pings from curl Pod to ClusterIP service	15
Figure 23 - Ping from one replica pod to another	15
Figure 24 - Ping from one replica pod to the ClusterIP service	16
Figure 25 - Pods in the kube-system namespace	16
Figure 26 - DNS ClusterIP service.....	16
Figure 27 - IP address of the DNS server on multiple Pods	17
Figure 28 - NodePort Service	17
Figure 29 - Replica Pods rescaled	18
Figure 30 - Webterm browser HTTP request.....	18
Figure 31 - Consecutive HTTP requests to UB2	19
Figure 32 - HTTP request to UB2 assigned to Pod at UB3	19
Figure 33 - HTTP message flow Request to UB2.....	20
Figure 34 - Docker swarm architecture	21
Figure 35 - Docker services in Docker Swarm.....	21
Figure 36 - Load Balancer deployment	22
Figure 37 - Load balancing container configuration file	22

Figure 38 - HTTP GET Request from Webterm browser	22
Figure 40 - Containers in ObiWan	23
Figure 39 - Containers in ObiTwo	23
Figure 41 - Successive HTTP requests from Toolbox	23

1. Kubernetes networking

The aim of this section is to explore Kubernetes and Kubernetes networking. The laboratory topics analysed are:

- I. Kubernetes control plane
- II. Deploying a single pod
- III. Deploying a single pod through a YAML manifest
- IV. Deploying a ReplicaSet
- V. Deploying a ClusterIP service
- VI. Kubernetes DNS
- VII. Deploying a NodePort service

Network Topology:

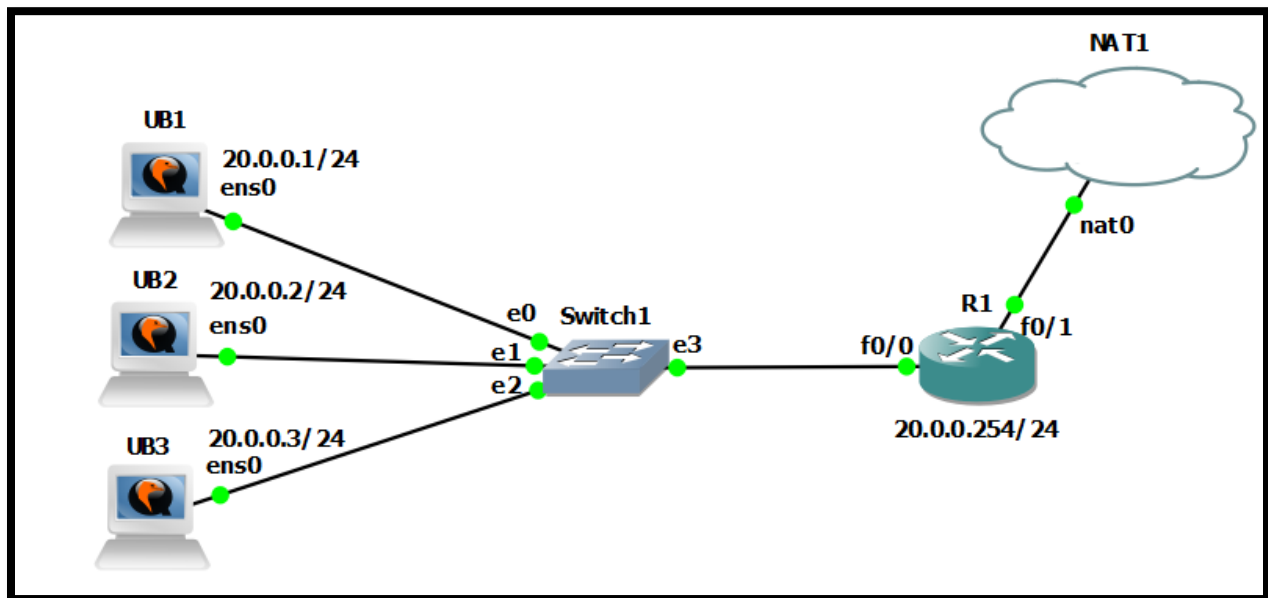


Figure 1 - Kubernetes networking network topology

Throughout this first part of the project, the study will focus on the network presented in Figure 1. Applied in this architecture are three Ubuntu 18.04.3 VMs, simulating Kubernetes nodes, one switch and one Cisco 3725 Router (R1).

The setup of the Ubuntu servers consisted of the installation of **Docker** and **Kubeadm**, the configuration of each machine's interfaces, and a **Kubernetes cluster**. These configurations are provided in [Annex A](#).

1.1. Kubernetes control plane

A Kubernetes cluster is composed of a set of worker machines, called nodes, that run containerised applications. These worker nodes host the Pods, which are the components of the application workload.

This study starts by analysing the exchanged traffic between the worker nodes and the master node. (Kubernetes, 2022)

The Control Plane:

In Kubernetes, the control plane manages the cluster's worker nodes and Pods. In our topology and in production environments, the control plane runs across multiple machines, providing fault tolerance and high availability.

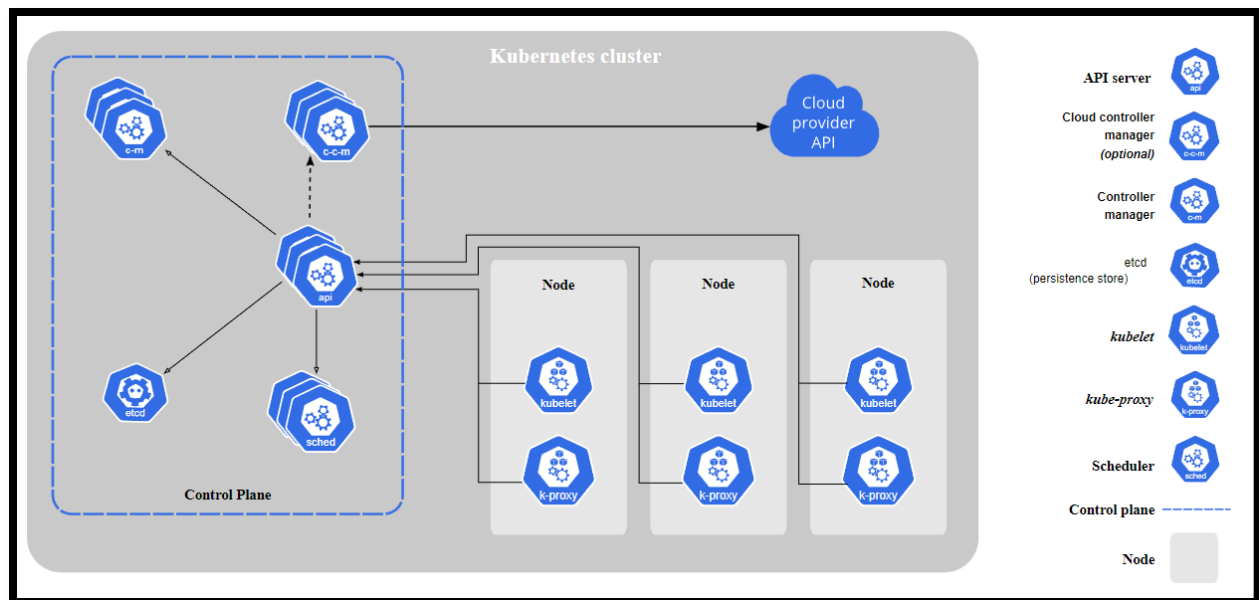


Figure 2 - Kubernetes cluster components (Kubernetes, 2022)

When running Kubernetes in an environment with strict network boundaries, it is helpful to be aware of the ports and protocols used by Kubernetes components. Table 1 shows the default ports and protocols used by the control plane with the reminder that all default ports can be overridden. (Kubernetes, 2022)

Table 1 - Control plane Ports and Protocols (Kubernetes, 2022)

Protocol	Direction	Port Range	Service	Used By
TCP	Inbound	6443	Kubernetes API server	All
TCP	Inbound	2379-2380	etcd server client API	kube-apiserver, etcd
TCP	Inbound	10250	Kubelet API	Self, Control plane
TCP	Inbound	10259	kube-scheduler	Self
TCP	Inbound	10257	kube-controller-manager	Self

In the following Figure 3 and Figure 4, it is possible to confirm two packets between the master and the worker nodes using port 6443 with Kubernetes API server information.

The packets from the master to the worker nodes use TLS1.2, while the packets from the worker nodes to the master use TCP.

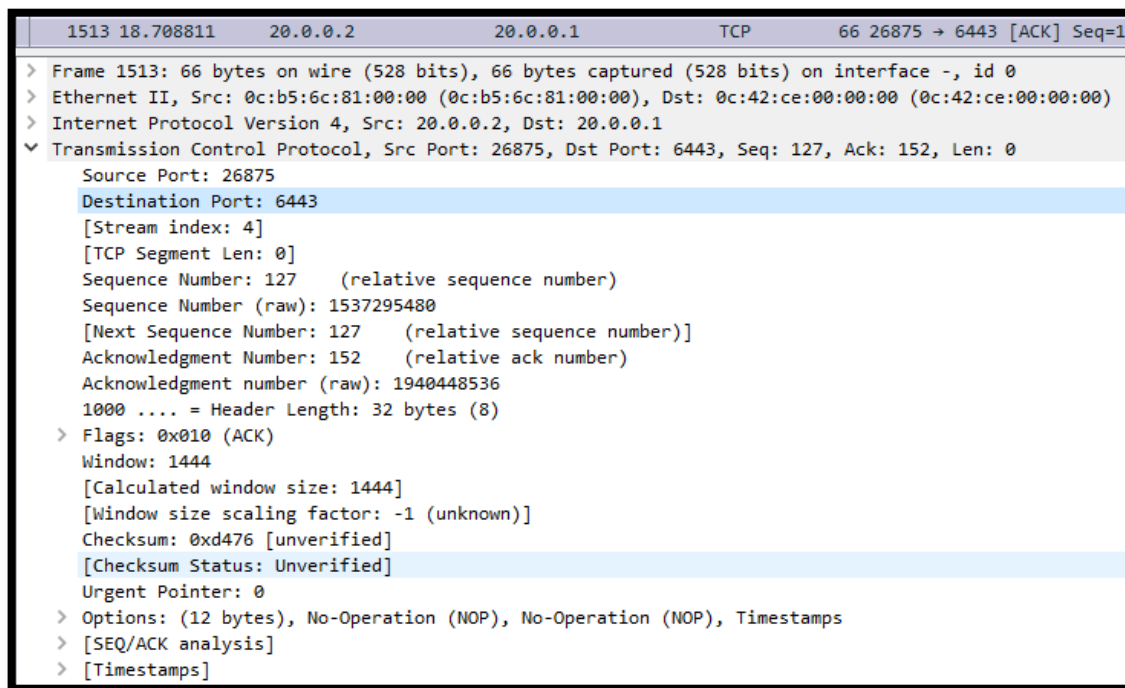


Figure 3 - Control plane packet from worker (UB1) to the master (UB2)

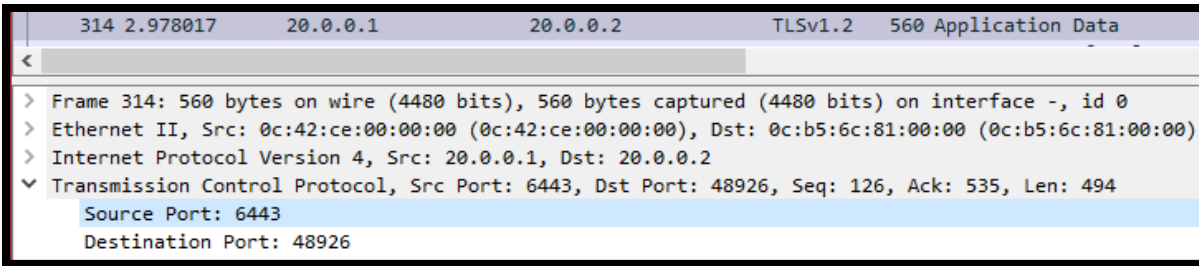


Figure 4 - Control plane packet from master (UB1) to the worker (UB2)

1.2. Deploying a single pod

Pods are the smallest deployable units of computing that can be created and managed in Kubernetes. A Pod is a group of one or more containers with shared storage and network resources and a specification for how to run the containers.

In terms of Docker concepts, a Pod is similar to a group of Docker containers with shared namespaces and shared filesystem volumes. (Kubernetes, 2022)

This section presents the deployment of a single pod with an nginx container. The complete configuration can be found in [Annex A](#).

Proof of Concept:

1. We create the pod named **mypod** with a nginx image. The pod is initially in a ContainerCreating state and then changes to a Running state. The created pod can be seen Running in Figure 5, with the IP address of 10.32.0.2 and located in node ub3.

```
root@UB1:/home/gns3# kubectl get pods -o wide
NAME      READY   STATUS    RESTARTS   AGE   IP        NODE   NOMINATED NODE   READINESS GATES
mypod     1/1     Running   1 (2m52s ago)  50m   10.32.0.2  ub3    <none>           <none>
root@UB1:/home/gns3#
```

Figure 5 - mypod IP address and node

2. To explore the Pod, a shell is created. In Figure 6, that same shell can be observed along with some commands demonstrating some of the Pod's information.

```
root@UB1:/home/gns3# kubectl exec -it mypod -- /bin/bash
root@mypod:/# ls
bin  docker-entrypoint.d  home  media  proc  sbin  tmp
boot  docker-entrypoint.sh  lib   mnt    root  srv   usr
dev   etc                  lib64  opt    run   sys   var
root@mypod:/# cat /etc/hosts
# Kubernetes-managed hosts file.
127.0.0.1    localhost
::1         localhost ip6-localhost ip6-loopback
fe00::0     ip6-localnet
fe00::0     ip6-mcastprefix
fe00::1     ip6-allnodes
fe00::2     ip6-allrouters
10.32.0.2    mypod
root@mypod:/# cat /etc/hostname
mypod
root@mypod:/# cat /usr/share/nginx/html/index.html
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
```

Figure 6 - Inside mypod shell

3. To prove that the Pods' service is operational, curl is installed, and a GET request is sent to the nginx server. Figure 7 shows a successful message from that same request.

```
root@mypod:/# curl http://localhost/
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
```

Figure 7 - GET request to the nginx server from mypod

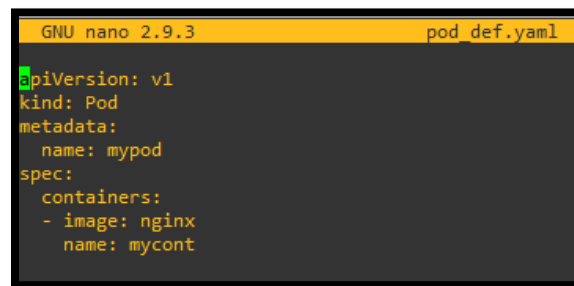
1.3. Deploying a single pod through a YAML manifest

The most common way of deploying a Pod is through a YAML manifest. A manifest specifies the desired state of an object that Kubernetes will maintain when that same manifest is applied. Each configuration file can contain multiple manifests.

This section presents the deployment of the same pod of the previous section, [1.2](#), but using a YAML manifest.

Proof of Concept:

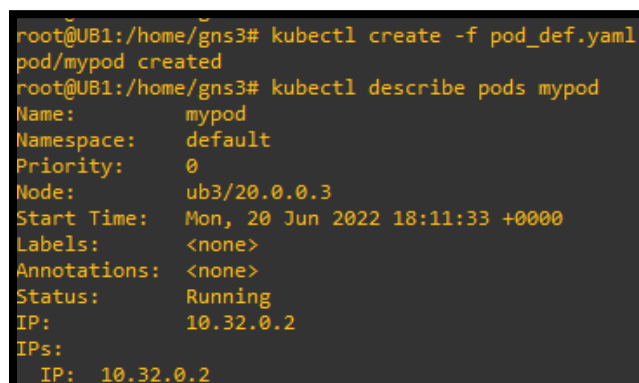
1. This procedure starts with the creation of a .yaml file specifying all the Pod's specifications. The contents of the file are shown in Figure 8.



```
GNU nano 2.9.3 pod_def.yaml
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
  - image: nginx
    name: mycont
```

Figure 8 - mypod YAML Manifest

2. Now the command **kubectl create -f pod_def.yaml** is executed and the Pod is created as shown in Figure 9.



```
root@UB1:/home/gns3# kubectl create -f pod_def.yaml
pod/mypod created
root@UB1:/home/gns3# kubectl describe pods mypod
Name:          mypod
Namespace:     default
Priority:       0
Node:          ub3/20.0.0.3
Start Time:    Mon, 20 Jun 2022 18:11:33 +0000
Labels:        <none>
Annotations:   <none>
Status:        Running
IP:            10.32.0.2
IPs:
  IP: 10.32.0.2
```

Figure 9 - Pod creation and Information

3. If the same command is executed twice, an error is displayed. So, it can be concluded that running the Pod creation command several times is not the correct way to generate replicas of a specific Pod, as seen in Figure 10.

```
root@UB1:/home/gns3# kubectl create -f pod_def.yaml
Error from server (AlreadyExists): error when creating "pod_def.yaml": pods "mypod" already exists
root@UB1:/home/gns3#
```

Figure 10 - Repeating the creation command

- The choice of the node where the Pod is created has been done automatically so far. To manually deploy the Pod in a specific node a new entry called **nodeName** must be added under **specs** on the YAML manifest. The Pod is now created again but now with the specification that it should be created on node UB2, as shown in Figure 11 and Figure 12.

```
GNU nano 2.9.3 pod_def.yaml
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  nodeName: ub2
  containers:
  - image: nginx
    name: mycont
```

Figure 11 - Modified YAML Manifest with manual node choice

```
root@UB1:/home/gns3# kubectl get pods -o wide
NAME    READY   STATUS    RESTARTS   AGE   IP          NODE   NOMINATED NODE   READINESS GATES
mypod   1/1     Running   0           43s   10.38.0.1   ub2    <none>           <none>
root@UB1:/home/gns3#
```

Figure 12 - mypod created in node UB2

1.4. Deploying a ReplicaSet

A ReplicaSet aims to maintain a stable set of replica Pods running at a given time.

A ReplicaSet is defined with fields including a selector that specifies how to identify Pods, several replicas and a Pod template specifying the data of new Pods it should create to meet the number of replicas required. (Kubernetes, 2022)

This section demonstrates the deployment of five replicas of an nginx pod using a YAML manifest of a ReplicaSet. The complete configuration can be found in [Annex A](#).

Proof of Concept:

- The generation of five replicas of a nginx pod will be done through a YAML manifest of a ReplicaSet. The respective file can be seen in Figure 13.

```

GNU nano 2.9.3 rep_def1.yaml
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myrep
spec:
  replicas: 5
  selector:
    matchLabels:
      app: mynginx
  template:
    metadata:
      labels:
        app: mynginx
    spec:
      containers:
      - name: mycont
        image: nginx

```

Figure 13 - ReplicaSet YAML Manifest

2. After the creation of the manifest, the set of replicas is created. The result and the main characteristics of the ReplicaSet can be seen in Figure 14, and the IP addresses plus the assigned nodes of each replica in Figure 15.

```

root@UB1:/home/gns3# kubectl create -f rep_def1.yaml
replicaset.apps/myrep created
root@UB1:/home/gns3# kubectl get replicaset -o wide
NAME      DESIRED  CURRENT  READY  AGE   CONTAINERS  IMAGES  SELECTOR
myrep     5         5         0      13s   mycont      nginx   app=mynginx

```

Figure 14 - ReplicaSet creation and Pods characteristics

```

root@UB1:/home/gns3# kubectl get pods -o wide
NAME      READY  STATUS   RESTARTS  AGE   IP            NODE    NOMINATED NODE  READINESS GATES
myrep-9bk2f  1/1    Running   0          87s   10.38.0.3     ub2     <none>           <none>
myrep-hc76x  1/1    Running   0          87s   10.32.0.3     ub3     <none>           <none>
myrep-jhc28  1/1    Running   0          87s   10.38.0.2     ub2     <none>           <none>
myrep-lsmz7  1/1    Running   0          87s   10.38.0.1     ub2     <none>           <none>
myrep-xphp8  1/1    Running   0          87s   10.32.0.2     ub3     <none>           <none>
root@UB1:/home/gns3#

```

Figure 15 - Pod's IP addresses and assigned nodes

3. Unlike when a user directly creates Pods, a ReplicaSet replaces Pods that are deleted or terminated for any reason, such as node failure or disruptive node maintenance, as a kernel upgrade. For this reason, it is recommended to use a ReplicaSet even if the application requires only a single Pod. Figure 16 shows the deletion of a Pod from the ReplicaSet that is immediately recreated in a different node, so it can be concluded that the pod that replaces the deleted is not always placed in the same node.

```

root@UB1:/home/gns3# kubectl delete pods myrep-9bk2f
pod "myrep-9bk2f" deleted
root@UB1:/home/gns3# kubectl get pods -o wide

```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
myrep-hc76x	1/1	Running	0	3m1s	10.32.0.3	ub3	<none>		<none>	
myrep-jhc28	1/1	Running	0	3m1s	10.38.0.2	ub2	<none>		<none>	
myrep-lsmz7	1/1	Running	0	3m1s	10.38.0.1	ub2	<none>		<none>	
myrep-n5hvj	0/1	ContainerCreating	0	6s	<none>	ub3	<none>		<none>	
myrep-xphp8	1/1	Running	0	3m1s	10.32.0.2	ub3	<none>		<none>	

```

root@UB1:/home/gns3#

```

Figure 16 - Pod deletion and replacement

- Scaling a ReplicaSet to create more replicas is very straightforward. Using the command **kubectl scale replicaset myrep --replicas=7** two more replicas are immediately created, as shown in Figure 17.

```

root@UB1:/home/gns3# kubectl get pods -o wide

```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
myrep-4vxhs	0/1	ContainerCreating	0	3s	<none>	ub2	<none>		<none>	
myrep-bkfcc	0/1	ContainerCreating	0	3s	<none>	ub2	<none>		<none>	
myrep-hc76x	1/1	Running	0	4m10s	10.32.0.3	ub3	<none>		<none>	
myrep-jhc28	1/1	Running	0	4m10s	10.38.0.2	ub2	<none>		<none>	
myrep-lsmz7	1/1	Running	0	4m10s	10.38.0.1	ub2	<none>		<none>	
myrep-n5hvj	1/1	Running	0	75s	10.32.0.4	ub3	<none>		<none>	
myrep-xphp8	1/1	Running	0	4m10s	10.32.0.2	ub3	<none>		<none>	

```

root@UB1:/home/gns3#

```

Figure 17 - Scaling the ReplicaSet to seven Pods

- When managing a ReplicaSet, there is no need to delete one Pod at a time. As shown in Figure 18, one can delete all Pods associated with a specific ReplicaSet.

```

root@UB1:/home/gns3# kubectl delete replicaset myrep
replicaset.apps "myrep" deleted
root@UB1:/home/gns3# kubectl get pods -o wide

```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
myrep-4vxhs	1/1	Terminating	0	47s	10.38.0.4	ub2	<none>		<none>	
myrep-bkfcc	1/1	Terminating	0	47s	10.38.0.3	ub2	<none>		<none>	
myrep-hc76x	1/1	Terminating	0	4m54s	10.32.0.3	ub3	<none>		<none>	
myrep-jhc28	1/1	Terminating	0	4m54s	10.38.0.2	ub2	<none>		<none>	
myrep-lsmz7	1/1	Terminating	0	4m54s	10.38.0.1	ub2	<none>		<none>	
myrep-n5hvj	1/1	Terminating	0	119s	10.32.0.4	ub3	<none>		<none>	
myrep-xphp8	1/1	Terminating	0	4m54s	10.32.0.2	ub3	<none>		<none>	

```

root@UB1:/home/gns3#

```

Figure 18 - Deleting the entire ReplicaSet

1.5. Deploying a ClusterIP service

Kubernetes provide Pods with their own IP addresses and a single DNS name for a set of Pods and can load-balance across them.

In Kubernetes, a Service is an abstraction that defines a logical set of Pods and a policy by which to access them. This kind of pattern is also sometimes used to apply a microservice.

There are several types of services in Kubernetes, the one applied in this exercise is called a **ClusterIP service**. This is a default service type, often used in Kubernetes. (Kubernetes , 2022)

Imagine deploying a microservices application in a cluster where a Pod runs a microservice container. That Pod will use a determined port number for that container and get an IP address from a Worker Node's range of IP addresses. With that port number and IP address the Pod can provide access to the microservice container. When creating a ReplicaSet for that specific microservice container the replica Pod opens the same port number but with a different IP address, just as demonstrated in Figure 19. The IP address of this replica may belong to a different range of IP addresses if it is created in a different worker node.

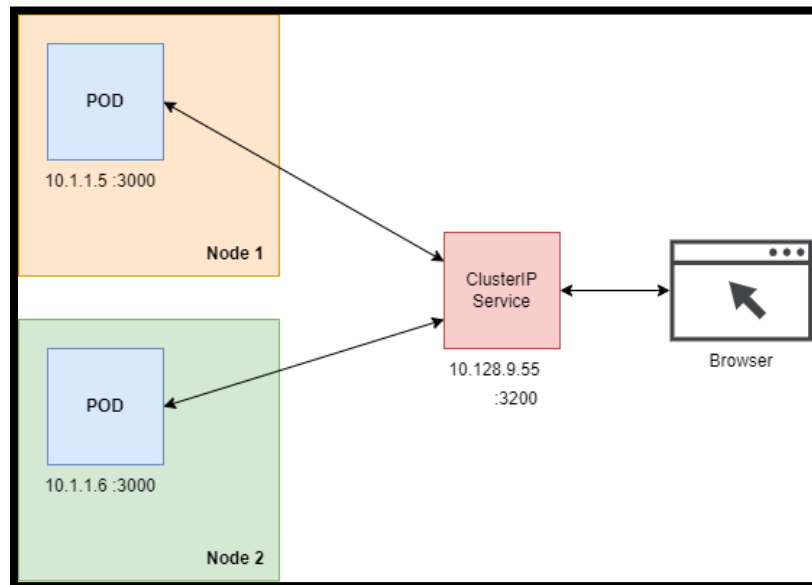


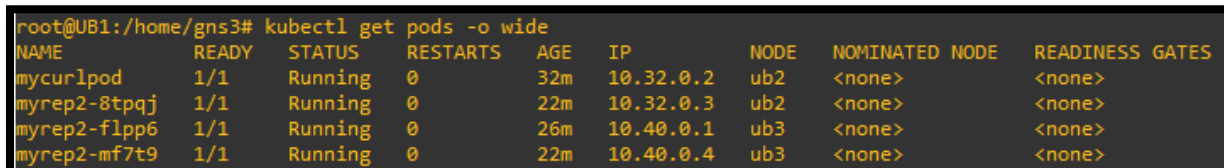
Figure 19 - ClusterIP service diagram

In the above situation, a ClusterIP service abstracts the ReplicaSet created for the microservice. When requests are made to that specific microservices, for example through a browser, a ClusterIP service with a certain port number and IP address handles those requests through the replicas. (Nana, 2020)

Proof of Concept:

This laboratory exercise demonstrates the deployment of a ClusterIP Service using a simple *nginxdemos/hello* application and a ReplicaSet. The complete configuration of this lab can be found in [Annex A](#).

1. The first step is to create a Pod with curl installed in UB1. This Pod will be used to test the load balancing service. To generate this Pod, we create a Pod with a simple nginx image and then enter the Pod to install curl.
2. Then a ReplicaSet is created with three *nginxdemos/hello* image Pods. The replicas and the Pod assembled in the previous step can be seen in Figure 20.

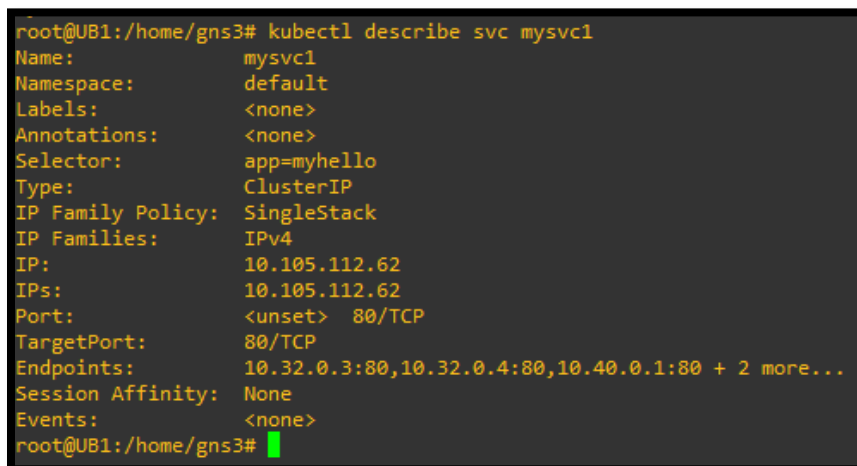


```
root@UB1:/home/gns3# kubectl get pods -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE	READINESS	GATES
mycurlpod	1/1	Running	0	32m	10.32.0.2	ub2	<none>		<none>	
myrep2-8tpqj	1/1	Running	0	22m	10.32.0.3	ub2	<none>		<none>	
myrep2-flpp6	1/1	Running	0	26m	10.40.0.1	ub3	<none>		<none>	
myrep2-mf7t9	1/1	Running	0	22m	10.40.0.4	ub3	<none>		<none>	

Figure 20 - Replica and Curl Pods

3. Then follows the deployment of the ClusterIP service, where it is specified that the pods containing the **app: hello** will be handled by the ClusterIP. Now with the ReplicaSet and ClusterIP service created, we scale the ReplicaSet from three to five replicas. The information regarding the created service can be found in Figure 21, which presents the service's IP address, Selector, the number of Endpoints and more.



```
root@UB1:/home/gns3# kubectl describe svc mysvc1
```

Name:	mysvc1
Namespace:	default
Labels:	<none>
Annotations:	<none>
Selector:	app=myhello
Type:	ClusterIP
IP Family Policy:	SingleStack
IP Families:	IPv4
IP:	10.105.112.62
IPs:	10.105.112.62
Port:	<unset> 80/TCP
TargetPort:	80/TCP
Endpoints:	10.32.0.3:80,10.32.0.4:80,10.40.0.1:80 + 2 more...
Session Affinity:	None
Events:	<none>

Figure 21 - ClusterIP service description

4. Using the command **curl -s http://10.105.112.62 | grep name** from inside the curl equipped pod created earlier, we receive a response indicating the name of the pod that was contacted by curl. Repeating this process reveals no patterns; therefore, the load balancing algorithm is random, as illustrated in Figure 22.

```

root@UB1:/home/gns3# kubectl get services -o wide
NAME          TYPE        CLUSTER-IP    EXTERNAL-IP    PORT(S)    AGE    SELECTOR
kubernetes    ClusterIP   10.96.0.1     <none>         443/TCP    27h    <none>
mysvc1        ClusterIP   10.105.112.62 <none>         80/TCP     5m53s  app=myhello

root@UB1:/home/gns3# kubectl exec -it mycurlpod -- /bin/bash
root@mycurlpod:/# curl -s http://10.105.112.62 | grep name
<p><span>Server&nbsp;name:</span> <span>myrep2-8tpqj</span></p>
root@mycurlpod:/# curl -s http://10.105.112.62 | grep name
<p><span>Server&nbsp;name:</span> <span>myrep2-flpp6</span></p>
root@mycurlpod:/# curl -s http://10.105.112.62 | grep name
<p><span>Server&nbsp;name:</span> <span>myrep2-flpp6</span></p>
root@mycurlpod:/# curl -s http://10.105.112.62 | grep name
<p><span>Server&nbsp;name:</span> <span>myrep2-8tpqj</span></p>
root@mycurlpod:/# curl -s http://10.105.112.62 | grep name
<p><span>Server&nbsp;name:</span> <span>myrep2-flpp6</span></p>
root@mycurlpod:/# curl -s http://10.105.112.62 | grep name
<p><span>Server&nbsp;name:</span> <span>myrep2-qg6xf</span></p>
root@mycurlpod:/# curl -s http://10.105.112.62 | grep name
<p><span>Server&nbsp;name:</span> <span>myrep2-8tpqj</span></p>
root@mycurlpod:/# curl -s http://10.105.112.62 | grep name
<p><span>Server&nbsp;name:</span> <span>myrep2-flpp6</span></p>
root@mycurlpod:/# curl -s http://10.105.112.62 | grep name
<p><span>Server&nbsp;name:</span> <span>myrep2-mf7t9</span></p>
root@mycurlpod:/# curl -s http://10.105.112.62 | grep name
<p><span>Server&nbsp;name:</span> <span>myrep2-qg6xf</span></p>
root@mycurlpod:/# curl -s http://10.105.112.62 | grep name
<p><span>Server&nbsp;name:</span> <span>myrep2-8tpqj</span></p>
root@mycurlpod:/# exit
exit
root@UB1:/home/gns3#

```

Figure 22 - Pings from curl Pod to ClusterIP service

- Now from one replica Pod, we ping another replica pod, and we can see in Figure 23 that the ping is successful. However, the same is not observed, in Figure 24, when we try to ping the ClusterIP service; this is because the Cluster IP is an abstraction of the service provided by the pods and does not reply to requests; it merely forwards them to a randomly chosen pod.

```

root@UB1:/home/gns3# kubectl get pods -o wide
NAME          READY    STATUS    RESTARTS    AGE    IP           NODE    NOMINATED NODE    READINESS GATES
mycurlpod     1/1      Running   0            19m    10.32.0.2    ub2     <none>             <none>
myrep2-8tpqj  1/1      Running   0            9m52s  10.32.0.3    ub2     <none>             <none>
myrep2-flpp6  1/1      Running   0            13m    10.40.0.1    ub3     <none>             <none>
myrep2-mf7t9  1/1      Running   0            9m52s  10.40.0.4    ub3     <none>             <none>
myrep2-qg6xf  1/1      Running   0            10m    10.32.0.4    ub2     <none>             <none>
myrep2-z4mhq  1/1      Running   0            13m    10.40.0.3    ub3     <none>             <none>

root@UB1:/home/gns3# kubectl exec myrep2-8tpqj ping 10.40.0.3
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in a future version. Use kubectl exec [POD] -- [COMMAND].
PING 10.40.0.3 (10.40.0.3): 56 data bytes
64 bytes from 10.40.0.3: seq=0 ttl=64 time=5.682 ms
64 bytes from 10.40.0.3: seq=1 ttl=64 time=0.668 ms
64 bytes from 10.40.0.3: seq=2 ttl=64 time=0.758 ms
64 bytes from 10.40.0.3: seq=3 ttl=64 time=0.684 ms
64 bytes from 10.40.0.3: seq=4 ttl=64 time=0.977 ms
64 bytes from 10.40.0.3: seq=5 ttl=64 time=0.915 ms
^C
root@UB1:/home/gns3#

```

Figure 23 - Ping from one replica pod to another


```

root@UB1:/home/gns3# kubectl exec myrep2-8tpqj ping 10.105.112.62
kubectl exec [POD] [COMMAND] is DEPRECATED and will be removed in
ead.
PING 10.105.112.62 (10.105.112.62): 56 data bytes
^C
root@UB1:/home/gns3#

```

Figure 24 - Ping from one replica pod to the ClusterIP service

1.6. Kubernetes DNS

In Kubernetes the DNS system schedules a DNS Pod and Service on the cluster and configures the kubelets to inform individual containers to use the DNS Service's IP to resolve DNS names. (Kubernetes , 2022)

Every Service Defined in the cluster is assigned a DNS name, including the DNS server itself. The complete configuration can be consulted in [Annex A](#).

Proof of Concept:

1. To learn more about how DNS works in a Kubernetes cluster, we start by analysing the *kube-system* namespace. In Figure 25, we can see two DNS servers with IP addresses: 10.40.0.2 and 10.38.0.1. Figure 26 shows that the DNS servers are gathered in a ClusterIP service named **kube-dns**.

```

root@UB1:/home/gns3# kubectl get pods -n kube-system -o wide
NAME                                READY  STATUS   RESTARTS   AGE   IP            NODE   NOMINATED NODE   READINESS GATES
coredns-6d4b75cb6d-f8ch1            1/1    Running  1 (110m ago)  120m  10.40.0.2     ub3    <none>            <none>
coredns-6d4b75cb6d-vsw84            1/1    Running  3 (96m ago)   120m  10.38.0.1     ub1    <none>            <none>
etcd-ub1                             1/1    Running  14 (96m ago)  28h   20.0.0.1     ub1    <none>            <none>
kube-apiserver-ub1                  1/1    Running  13 (96m ago)  28h   20.0.0.1     ub1    <none>            <none>
kube-controller-manager-ub1         1/1    Running  9 (96m ago)   28h   20.0.0.1     ub1    <none>            <none>
kube-proxy-9864s                    1/1    Running  4 (110m ago)  28h   20.0.0.3     ub3    <none>            <none>
kube-proxy-hvpmr                     1/1    Running  6 (96m ago)   28h   20.0.0.1     ub1    <none>            <none>
kube-proxy-ksp5m                     1/1    Running  3 (110m ago)  28h   20.0.0.2     ub2    <none>            <none>
kube-scheduler-ub1                  1/1    Running  9 (96m ago)   28h   20.0.0.1     ub1    <none>            <none>
weave-net-26gmr                     2/2    Running  7 (110m ago)  28h   20.0.0.2     ub2    <none>            <none>
weave-net-5ljzk                     2/2    Running  13 (95m ago)  28h   20.0.0.1     ub1    <none>            <none>
weave-net-n4lzk                     2/2    Running  9 (110m ago)  28h   20.0.0.3     ub3    <none>            <none>
root@UB1:/home/gns3#

```

Figure 25 - Pods in the kube-system namespace

```

root@UB1:/home/gns3# kubectl get service -n kube-system
NAME      TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
kube-dns  ClusterIP   10.96.0.10   <none>        53/UDP,53/TCP,9153/TCP  28h
root@UB1:/home/gns3#

```

Figure 26 - DNS ClusterIP service

2. By consulting the *resolv.conf* file in several Pods, we can conclude that no Pod uses the IP address of a DNS server directly; they instead use the IP address of the ClusterIP service, **kube-dns**, that handles the DNS requests for the available DNS servers. Such a search can be verified in Figure 27.

```

root@UB1:/home/gns3# kubectl exec myrep2-8tpqj -- cat /etc/resolv.conf
search default.svc.cluster.local svc.cluster.local cluster.local
nameserver 10.96.0.10
options ndots:5
root@UB1:/home/gns3# kubectl exec myrep2-flpp6 -- cat /etc/resolv.conf
search default.svc.cluster.local svc.cluster.local cluster.local
nameserver 10.96.0.10
options ndots:5
root@UB1:/home/gns3# kubectl exec myrep2-z4mhq -- cat /etc/resolv.conf
search default.svc.cluster.local svc.cluster.local cluster.local
nameserver 10.96.0.10
options ndots:5
root@UB1:/home/gns3#

```

Figure 27 - IP address of the DNS server on multiple Pods

3. To test this DNS service, we access the service mysvc1 using the service's DNS name. Using the commands presented below, it is possible to conclude that all these different commands work because they represent a smaller section of the search parameter shown in Figure 27 defined in the DNS server. As previously confirmed, all Pods use it.

```

1. #Access the service mysvc1 using the service's DNS name
2. kubectl exec mycurlpod -- curl -s http://mysvc1
3. kubectl exec mycurlpod -- curl -s http://kubectl exec mycurlpod -- curl -s
  http://mysvc1.kube-system
4. kubectl exec mycurlpod -- curl -s http://kubectl exec mycurlpod -- curl -s
  http://mysvc1.kube-system.svc
5. kubectl exec mycurlpod -- curl -s http://kubectl exec mycurlpod -- curl -s
  http://mysvc1.kube-system.svc.cluster.local

```

1.7. Deploying a NodePort service

The NodePort service allows the connection of external users to the cluster.

When creating a Service of type NodePort, Kubernetes provides a **nodePort** value. Then the Service is accessible by using the IP address of any node along with the **nodePort** value. (Google, 2022)

Proof of Concept:

1. A NodePort service is created using a YAML file selecting an app: myhello service and a nodePort: 30123. The service's IP address and port number can be seen in Figure 28.

```

root@UB1:/home/gns3# kubectl get svc -o wide
NAME          TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE   SELECTOR
kubernetes    ClusterIP   10.96.0.1    <none>        443/TCP          45h   <none>
mysvc2        NodePort    10.96.1.251  <none>        80:30123/TCP     8s    app=myhello

```

Figure 28 - NodePort Service

- Then the ReplicaSet is scaled to three replicas, as we can see in Figure 29, and a webterm and Toolbox are connected to the network. Using the IP address of one of the worker nodes and the port number provided by the NodePort service, we can confirm the service's success in Figure 30.

```
root@UB1:/home/gns3# kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP            NODE   NOMINATED NODE   READINESS GATES
mycurlpod     1/1     Running   2 (36m ago)  17h   10.32.0.2     ub2    <none>            <none>
myrep2-flpp6   1/1     Running   2 (36m ago)  17h   10.40.0.1     ub3    <none>            <none>
myrep2-qg6xf   1/1     Running   2 (36m ago)  17h   10.32.0.4     ub2    <none>            <none>
myrep2-z4mhq   1/1     Running   2 (36m ago)  17h   10.40.0.3     ub3    <none>            <none>
```

Figure 29 - Replica Pods rescaled

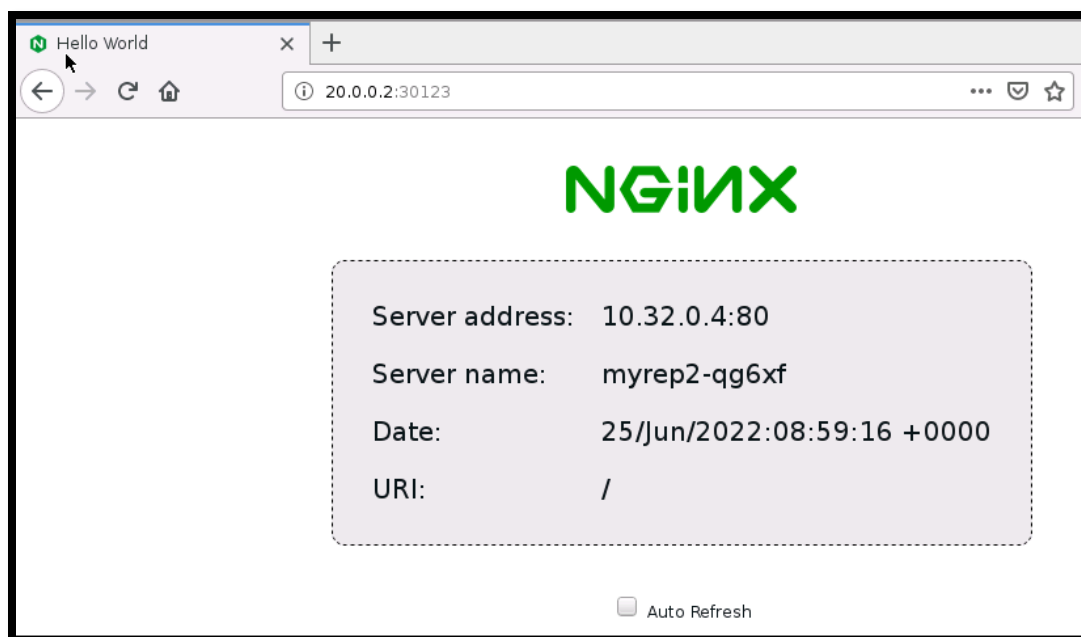


Figure 30 - Webterm browser HTTP request

- By using the Toolbox to access the containers, we can grep the name of the Pod used in each call. By doing this, it is possible to conclude that the type of load balancing algorithm is random since no pattern was identified in a consecutive experience, showed in Figure 31.

```

100 7217 0 7217 0 0 201k 0 ----- 207k
<p><span>Server<span>name:</span></span> <span>myrep2-flpp6</span></span></p>
root@Toolbox-1:~# curl http://20.0.0.2:30123 | grep name
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload  Total  Spent    Left  Speed
100 7217 0 7217 0 0 370k 0 ----- 370k
<p><span>Server<span>name:</span></span> <span>myrep2-qg6xf</span></span></p>
root@Toolbox-1:~# curl http://20.0.0.2:30123 | grep name
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload  Total  Spent    Left  Speed
100 7217 0 7217 0 0 306k 0 ----- 440k
<p><span>Server<span>name:</span></span> <span>myrep2-flpp6</span></span></p>
root@Toolbox-1:~# curl http://20.0.0.2:30123 | grep name
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload  Total  Spent    Left  Speed
100 7217 0 7217 0 0 1409k 0 ----- 1409k
<p><span>Server<span>name:</span></span> <span>myrep2-qg6xf</span></span></p>
root@Toolbox-1:~# curl http://20.0.0.2:30123 | grep name
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload  Total  Spent    Left  Speed
100 7217 0 7217 0 0 271k 0 ----- 271k
<p><span>Server<span>name:</span></span> <span>myrep2-z4mhq</span></span></p>
root@Toolbox-1:~# curl http://20.0.0.2:30123 | grep name
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload  Total  Spent    Left  Speed
100 7217 0 7217 0 0 783k 0 ----- 880k
<p><span>Server<span>name:</span></span> <span>myrep2-z4mhq</span></span></p>
root@Toolbox-1:~# curl http://20.0.0.2:30123 | grep name
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload  Total  Spent    Left  Speed
100 7217 0 7217 0 0 783k 0 ----- 1761k
<p><span>Server<span>name:</span></span> <span>myrep2-z4mhq</span></span></p>
root@Toolbox-1:~# curl http://20.0.0.2:30123 | grep name
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload  Total  Spent    Left  Speed
100 7217 0 7217 0 0 414k 0 ----- 414k
<p><span>Server<span>name:</span></span> <span>myrep2-z4mhq</span></span></p>
root@Toolbox-1:~#

```

Figure 31 - Consecutive HTTP requests to UB2

- By analysing the flow of HTTP messages in a Wireshark capture, it is possible to conclude that when a message arrives at a node but is assigned to a Pod on another node, the node contacted is the one that sends the HTTP response. Such an experiment is demonstrated in Figure 32, where an HTTP request is forwarded to UB2 but gets assigned to Pod myrep2-z4mhq located in UB3. In Figure 33, we can see that although assigned to a Pod on a different host, it is still UB2 that responds.

```

root@Toolbox-1:~# curl http://20.0.0.2:30123 | grep name
  % Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
                                 Dload  Upload  Total  Spent    Left  Speed
100 7217 0 7217 0 0 704k 0 ----- 704k
<p><span>Server<span>name:</span></span> <span>myrep2-z4mhq</span></span></p>
root@Toolbox-1:~#

```

Figure 32 - HTTP request to UB2 assigned to Pod at UB3

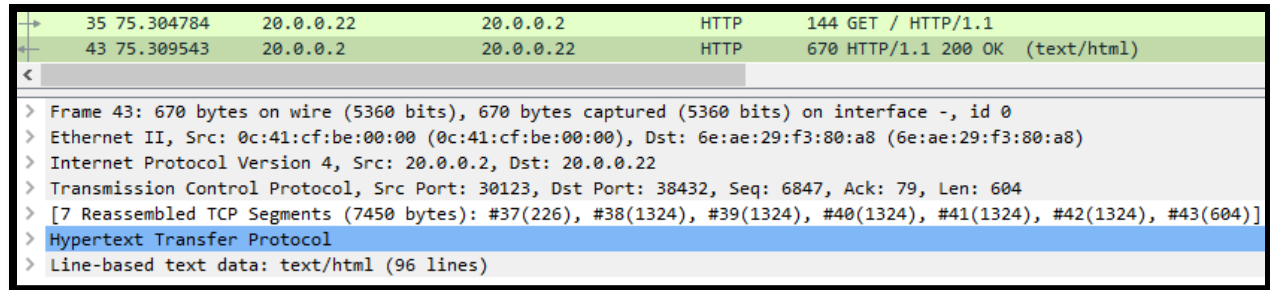


Figure 33 - HTTP message flow Request to UB2

2. Docker Swarm

This section aims to explore Docker technology, in particular comparing a service's deployment in swarm versus Kubernetes. The laboratory topics analysed are:

I. Deploying a Service in Docker Swarm

2.1. Deploying a Service in Docker Swarm

The objective of this exercise is to compare the deployment of services between a docker swarm and a Kubernetes cluster.

Docker swarm makes a cluster orchestration more manageable and allows the opportunity to publish ports for services. Usually, the only way to access containers is through their host machine's IP address, but in a swarm, all nodes participate in an ingress routing mesh.

Swarm mode Routing Mesh

Docker Engine swarm mode eases the process of publishing ports for services to make them available to resources outside the swarm. The routing mesh enables each node in the swarm to accept connections on published ports for any service in the swarm, even if no task is running on the node. (Docker, 2022)

The routing mesh routes all incoming requests to published ports on available nodes to an active container.

The ingress mesh uses the following ports:

- Port 7946 TCP/UDP for container network discovery.
- Port 4789 UDP for the container ingress network.

Network Architecture




Figure 34 - Docker swarm architecture

In this architecture, we use one docker swarm environment with ObiWan and ObiTwo and a separate Load Balancer machine; all are Ubuntu 18.04 machines. To connect them to the internet, we use NAT, a Cisco 3725 router and a standard Switch. The complete configuration can be found in [Annex A](#).

Proof of Concept:

To better understand this technology, we demonstrate the publication of a simple `nginxdemos/hello` web service using a Docker swarm and a Load Balancer.

1. First, we enable the docker swarm for two of the three nodes presented in Figure 34, ObiWan as the manager and ObiTwo as the worker. Then start the web host `nginxdemos/hello` service in two replicas and publish them to external port 8080, shown in Figure 35.

```

root@ObiWan:/home/gns3# docker service ls
ID            NAME      MODE     REPLICAS  IMAGE                PORTS
pp1rbsvejvcg backend replicated 2/2        nginxdemos/hello:latest *:8080->80/tcp

```

Figure 35 - Docker services in Docker Swarm

2. For consistency, the load balancer is deployed on a single node outside the swarm as shown in Figure 36. As all the published services are available through any of the swarm nodes thanks to the ingress routing, the load balancer can be set to use the swarm private IP addresses without concern about which node is hosting what service.

```
root@LB1:/home/gns3# docker node ls
ID                                HOSTNAME    STATUS    AVAILABILITY    MANAGER STATUS    ENGINE VERSION
xghfnklhf5crq5buk0qdmgqaa *     LB1        Ready    Active          Leader            20.10.7
```

Figure 36 - Load Balancer deployment

- Then we create the load balancing container using the configuration file presented in Figure 37. As we can see, the file contains the IP addresses of both swarm nodes, and it is specified to publish the container to port 80.

```
GNU nano 2.9.3 /data/loadbalancer/default.conf
server {
    listen 80;
    location / {
        proxy_pass http://backend;
    }
}
upstream backend {
    server 20.0.0.1:8080;
    server 20.0.0.2:8080;
}
```

Figure 37 - Load balancing container configuration file

- Now to test the load balancer we attach a Webterm to the network and execute an HTTP GET request from its web browser, using the IP of the load balancer and as shown in Figure 38 the connection to the web service is successful.




Figure 38 - HTTP GET Request from Webterm browser

- To better understand the type of load balancing method used in this solution we will execute multiple HTTP Get requests to the load balancer. For such a purpose we first, identify each container and where they are located; we can see the containers created in

node ObiWan and ObiTwo in Figure 40 and Figure 39, respectively. We can now conclude that the container used in the previous web browser request is located in ObiTwo.

```
root@ObiWan:/home/gns3# docker container ls
CONTAINER ID   IMAGE                                COM
3c93cc8ab44c   nginxdemos/hello:latest           "/c
461cc16whszga4xp8taa
928f3618c3ae   nginxdemos/hello:latest           "/c
yat0gg3qtjjaz11lefpc
root@ObiWan:/home/gns3#
```

Figure 40 - Containers in ObiWan

```
root@ObiTwo:/home/gns3# docker container ls
CONTAINER ID   IMAGE                                COMMA
8c37474c731e   nginxdemos/hello:latest           "/doc
ehvbpthgg2wpam1v9jmk
8c51ecfd239c   nginxdemos/hello:latest           "/doc
1ejd9k0s7loouaqcbjke
5c5400575769   nginxdemos/hello:latest           "/doc
hy3zqbmqz1aqd4q6ge0c
root@ObiTwo:/home/gns3#
```

Figure 39 - Containers in ObiTwo

- Analysing the successive requests presented in Figure 41, we can conclude that the load balancing algorithm is merely random with no pattern identified.

```
root@Toolbox-1:~# curl http://20.0.0.3 | grep name
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total     Spent    Left     Speed
100  7217    0  7217    0    0   281k      0 --:--:-- --:--:-- --:--:--   281k
<p><span>Server&nbsp;</span></span></p>
root@Toolbox-1:~# curl http://20.0.0.3 | grep name
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total     Spent    Left     Speed
100  7217    0  7217    0    0   640k      0 --:--:-- --:--:-- --:--:--   640k
<p><span>Server&nbsp;</span></span></p>
root@Toolbox-1:~# curl http://20.0.0.3 | grep name
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total     Spent    Left     Speed
100  7217    0  7217    0    0   281k      0 --:--:-- --:--:-- --:--:--   281k
<p><span>Server&nbsp;</span></span></p>
root@Toolbox-1:~# curl http://20.0.0.3 | grep name
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total     Spent    Left     Speed
100  7217    0  7217    0    0   440k      0 --:--:~ --:~:~ --:~:~   440k
<p><span>Server&nbsp;</span></span></p>
root@Toolbox-1:~# curl http://20.0.0.3 | grep name
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total     Spent    Left     Speed
100  7217    0  7217    0    0   640k      0 --:~:~ --:~:~ --:~:~   640k
<p><span>Server&nbsp;</span></span></p>
root@Toolbox-1:~# curl http://20.0.0.3 | grep name
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total     Spent    Left     Speed
100  7216    0  7216    0    0   880k      0 --:~:~ --:~:~ --:~:~   880k
<p><span>Server&nbsp;</span></span></p>
root@Toolbox-1:~#
```

Figure 41 - Successive HTTP requests from Toolbox

References:

- Docker. (2022). *Use swarm mode routing mesh*. Retrieved from Docker Documentation: <https://docs.docker.com/engine/swarm/ingress/>
- Google . (2022). *Services*. Retrieved from Google Cloud: <https://cloud.google.com/kubernetes-engine/docs/concepts/service>
- Kubernetes . (2022, May 30). *DNS for Services and Pods*. Retrieved from Kubernetes Documentation: <https://kubernetes.io/docs/concepts/services-networking/dns-pod-service/>
- Kubernetes . (2022, May 09). *Ports and Protocols*. Retrieved from Kubernetes Documentation: <https://kubernetes.io/docs/reference/ports-and-protocols/>
- Kubernetes . (2022, May 31). *Service*. Retrieved from Kubernetes Documentation: <https://kubernetes.io/docs/concepts/services-networking/service/>
- Kubernetes. (2022, April 30). *Kubernetes Components*. Retrieved from Kubernetes: <https://kubernetes.io/docs/concepts/overview/components/>
- Kubernetes. (2022, January 10). *Pods*. Retrieved from Kubernetes Documentation: <https://kubernetes.io/docs/concepts/workloads/pods/#:~:text=Pods%20are%20the%20smallest%20deployable,how%20to%20run%20the%20containers.>
- Kubernetes. (2022, June 07). *ReplicaSet*. Retrieved from Kubernetes Documentation: <https://kubernetes.io/docs/concepts/workloads/controllers/replicaset/>
- Nana, T. w. (2020, October 28). *Kubernetes Services explained | ClusterIP vs NodePort vs LoadBalancer vs Headless Service*. Retrieved from YouTube: <https://www.youtube.com/watch?v=T4Z7visMM4E>

Annex A: Configurations

Docker and Kubernetes installation in GNS3:

```

1. R1 Configuration:
2.
3. conf t
4. int f0/0
5. ip add 20.0.0.254 255.255.255.0
6. ip nat inside
7. no shut
8.
9. int f0/1
10. ip add dhcp
11. ip nat outside
12. no shut
13. exit
14. ip nat inside source list 1 interface FastEthernet0/1 overload
15. end
16.
17. conf t
18. access-list 1 permit 20.0.0.0 0.0.0.255
19. end
20.
21. #####
22. UB1 Configuration:
23.
24. hostnamectl set-hostname UB1
25. exec bash
26. nano /etc/netplan/50-cloud-init.yaml
27.
28. network:
29.   version: 2
30.   renderer: networkd
31.   ethernets:
32.     ens3:
33.       addresses:
34.         - 20.0.0.1/24
35.       gateway4: 20.0.0.254
36.       nameservers:
37.         addresses:
38.           - 8.8.8.8
39.
40. netplan apply
41. reboot
42.
43. #Docker installation
44. apt-get update
45. apt install docker.io
46. reboot
47. systemctl status docker
48. sudo apt-get update && sudo apt-get upgrade
49.
50. #Kubeadm installation
51. nano /etc/fstab
52. reboot
53.
54. nano /lib/systemd/system/docker.service
55. ExecStart=/usr/bin/dockerd --exec-opt native.cgroupdriver=systemd
56. systemctl status docker
57.
58. apt-get update

```

```

59. apt-get install -y apt-transport-https ca-certificates curl
60. curl -fsSLo /usr/share/keyrings/kubernetes-archive-keyring.gpg
   https://packages.cloud.google.com/apt/doc/apt-key.gpg
61. echo "deb [signed-by=/usr/share/keyrings/kubernetes-archive-keyring.gpg]
   https://apt.kubernetes.io/ kubernetes-xenial main" | sudo tee
   /etc/apt/sources.list.d/kubernetes.list
62.
63. apt-get update
64. apt-get install -y kubelet kubeadm kubectl
65. apt-mark hold kubelet kubeadm kubectl
66.
67. kubeadm config images pull
68.
69. #Docker images for labs
70. docker pull nginx
71. nano app.js
72. const http = require('http');
73. const os = require('os');
74. console.log("Kubia server starting...");
75. var handler = function(request, response) {
76.   console.log("Received request from " + request.connection.remoteAddress);
77.   response.writeHead(200);
78.   response.end("You've hit " + os.hostname() + "\n");
79. };
80. var www = http.createServer(handler);
81. www.listen(8080);
82.
83. nano Dockerfile
84. FROM node:7
85. ADD app.js /app.js
86. ENTRYPOINT ["node", "app.js"]
87.
88. docker build -t kubia .
89. docker image ls
90.
91. #Kubernetes cluster
92. kubeadm init --apiserver-advertise-address=20.0.0.1 --pod-network-cidr=10.32.0.0/12
93.
94. mkdir -p $HOME/.kube
95. cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
96. chown $(id -u):$(id -g) $HOME/.kube/config
97.
98. kubectl apply -f https://cloud.weave.works/k8s/net?k8s-version=$(kubectl version | base64 |
   tr -d '\n');
99.
100. kubectl get pods -A
101. kubectl get nodes
102. #####
103. Ubuntu 2 Configuration:
104.
105. hostnamectl set-hostname UB2
106. nano /etc/netplan/50-cloud-init.yaml
107.
108. network:
109.   version: 2
110.   renderer: networkd
111.   ethernets:
112.     ens3:
113.       addresses:
114.         - 20.0.0.2/24
115.       gateway4: 20.0.0.254
116.       nameservers:
117.         addresses:
118.           - 8.8.8.8
119. netplan apply

```

```

120. reboot
121.
122.
123. #Join worker
124. kubeadm join 20.0.0.1:6443 --token vjrxio.noa2wizoz107dsxu --discovery-token-ca-cert-hash
    sha256:b44227d48388f9d8c7edceacaa27a825141d7dd9055d642bbd9784c0371bf74f
125.
126. #####
127. Ubuntu 3 Configuration:
128.
129. hostnamectl set-hostname UB3
130. nano /etc/netplan/50-cloud-init.yaml
131.
132. network:
133.   version: 2
134.   renderer: networkd
135.   ethernets:
136.     ens3:
137.       addresses:
138.         - 20.0.0.3/24
139.       gateway4: 20.0.0.254
140.       nameservers:
141.         addresses:
142.           - 8.8.8.8
143.
144. netplan apply
145. reboot
146.
147. #Join worker
148. kubeadm join 20.0.0.1:6443 --token vjrxio.noa2wizoz107dsxu --discovery-token-ca-cert-hash
    sha256:b44227d48388f9d8c7edceacaa27a825141d7dd9055d642bbd9784c0371bf74f

```

Deploying a single pod

```

1. sudo /etc/init.d/networking restart
2. apt-get update
3. docker pull nginx
4.
5. kubectl get pods -A
6. kubectl run mypod --image=nginx
7. kubectl get pods
8.
9. #Identify the IP address assigned to the pod and the node the pod was assigned to
10. kubectl get pods -o wide
11. kubectl describe pods mypod * | grep Node * | grep IP
12.
13. #Explore the pod interior
14. kubectl exec -it mypod -- /bin/bash
15. ls
16. cat /etc/hosts
17. cat /etc/hostname
18. cat /usr/share/nginx/html/index.html
19.
20. #Install curl and send a GET request to the nginx server
21. apt-get update
22. apt-get install curl
23.
24. curl http://localhost/
25. exit

```

Deploying a ReplicaSet

```
1. #Create Replicas Manifest
2. nano rep_def1.yaml
3.
4. apiVersion: apps/v1
5. kind: ReplicaSet
6. metadata:
7.   name: myrep
8. spec:
9.   replicas: 5
10.  selector:
11.    matchLabels:
12.      app: mynginx
13.  template:
14.    metadata:
15.      labels:
16.        app: mynginx
17.    spec:
18.      containers:
19.      - name: mycont
20.        image: nginx
21.
22. #Create Replicas
23. kubectl create -f rep_def1.yaml
24. kubectl get replicaset -o wide
25. kubectl get pods -o wide
26.
27. #Delete one pod
28. kubectl delete pods myrep-9bk2f
29.
30. #Scale the ReplicaSet to seven pods
31. kubectl scale replicaset myrep --replicas=7
32. kubectl get pods -o wide
33.
34. #Delete the replica set
35. kubectl delete replicaset myrep
36.
```

Deploying a ClusterIP service

```
1. #Create a pod that runs curl:
2.
3. kubectl run mypod --image=nginx
4. kubectl get pods -A
5. kubectl exec -it mycurlpod -- /bin/bash
6. apt update
7. apt install curl
8. curl --version
9.
10. kubectl delete pod mycurlpod
11.
12. #Create the ReplicaSet
13. nano rep_def1.yaml
14.
15. apiVersion: apps/v1
16. kind: ReplicaSet
17. metadata:
18.   name: myrep2
19. spec:
20.   replicas: 3
21.   selector:
22.     matchLabels:
23.       app: myhello
24.   template:
25.     metadata:
26.       labels:
27.         app: myhello
28.     spec:
29.       containers:
30.       - name: mycont
31.         image: nginxdemos/hello
32.
33. kubectl create -f rep_def1.yaml
34.
35. kubectl get replicaset -o wide
36. kubectl get pods -A
37. kubectl get nodes
38.
39. kubectl delete replicaset myrep2
40.
41. #Deploy the ClusterIP service
42. nano svc_def1.yaml
43.
44. apiVersion: v1
45. kind: Service
46. metadata:
47.   name: mysvc1
48. spec:
49.   ports:
50.   - port: 80
51.     targetPort: 80
52.   selector:
53.     app: myhello
54.
55. kubectl create -f svc_def1.yaml
56.
57. kubectl get services -o wide
58. kubectl describe svc mysvc1
59. kubectl get endpoints mysvc1
60.
61. kubectl delete services mysvc1
```

```
62.
63. #Rescale the ReplicaSet to 5 pods
64. kubectl scale rs myrep2 --replicas=5
65. kubectl get pods -o wide
66.
67. #Type of load balancing algorithm is being used
68. kubectl get pods
69. kubectl exec -it mycurlpod -- /bin/bash
70. curl -s http://10.105.112.62 | grep name
71.
72. #Pings
73. kubectl get pods -o wide
74. kubectl exec myrep2-8tpqj ping 10.40.0.3
75. kubectl exec myrep2-8tpqj ping 10.105.112.62
```

Kubernetes DNS

```
1. #Analyze the pods of the kube-system namespace
2. kubectl get pods -n kube-system -o wide
3.
4. #Check that the DNS servers are gathered in a ClusterIP service
5. kubectl get service -n kube-system
6. kubectl get endpoints -n kube-system kube-dns
7.
8. #Look for the IP address of the DNS server
9. kubectl exec <pod name> -- cat /etc/resolv.conf
10.
11. #Access the service mysvc1 using the service's DNS name
12. kubectl exec mycurlpod -- curl -s http://mysvc1
13. kubectl exec mycurlpod -- curl -s http://kubectl exec mycurlpod -- curl -s
    http://mysvc1.kube-system
14. kubectl exec mycurlpod -- curl -s http://kubectl exec mycurlpod -- curl -s
    http://mysvc1.kube-system.svc
15. kubectl exec mycurlpod -- curl -s http://kubectl exec mycurlpod -- curl -s
    http://mysvc1.kube-system.svc.cluster.local
16.
17. #Delete the service
18. kubectl delete svc mysvc1
```

Deploying a Service in Docker Swarm

```
1. #Initiate Cluster manager
2. docker swarm init --advertise-addr 20.0.0.1
3.
4. #Join worker node
5. docker swarm join --token SWMTKN-1-1nku8t8e8g3bckl4fzu7hb2ovjcutg2jbl48qzfw9321tkh1uv-
   29z31vn96me8je9w7rug1n88j 20.0.0.1:2377
6.
7. #Create Replica service
8. docker service create --name backend --replicas 2 --publish 8080:80 nginxdemos/hello
9. docker service scale backend=5
10.
11. #Setup Load Balancer
12. docker swarm init --advertise-addr 20.0.0.3
13. mkdir -p /data/loadbalancer
14.
15. nano /data/loadbalancer/default.conf
16.
17. server {
18.     listen 80;
19.     location / {
20.         proxy_pass http://backend;
21.     }
22. }
23. upstream backend {
24.     server 20.0.0.1:8080;
25.     server 20.0.0.2:8080;
26. }
27.
28. docker service create --name loadbalancer --mount
   type=bind,source=/data/loadbalancer,target=/etc/nginx/conf.d --publish 80:80 nginx
29.
30. #Test Load balancer
31. http://20.0.0.3
32. curl http://<LoadBalancer IP address> | grep name
33.
```